

We Ain't Afraid of No File Fragmentation: Causes and Prevention of Its Performance Impact on Modern Flash SSDs

*Jun, Yuhun, Shinhyun Park, Jeong-Uk Kang, Sang-Hoon Kim,
and Euseong Seo.*

USENIX FAST'24

2024. 08. 14

Presented by Juhyun Kim & Yongmin Lee
(email) jhk@kiost.ac.kr , nascarf16@dankook.ac.kr

Contents

1. Fragmentation in HDD
2. Normal SSD
3. Fragmentation in SSD
4. Analysis of File Fragmentation
5. Approach
6. Evaluation
7. Conclusion

Fragmentation in HDD

- **Discontinue data blocks**

- Random access to scattered fragment

- Read performance bad !!!


- **Existing tool**

- Delay, pre allocation etc ...
 - But simultaneously multiple write or long time before additional file write

- Impossible avoid to fragmentation

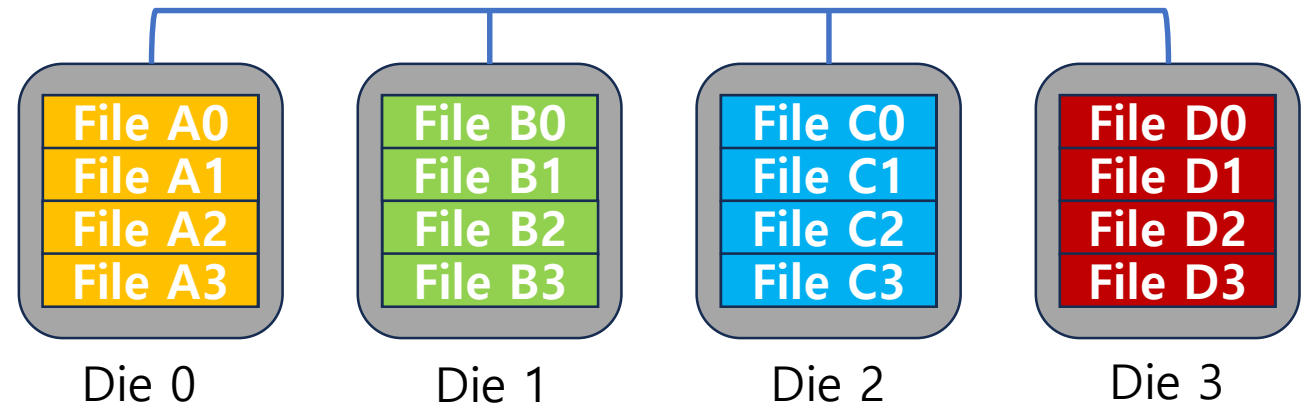

- **Main reason**

- Kernel I/O path, storage device interface, storage media access



Normal SSD

- SSD



Is performance always good?

Fragmentation in SSD

- SSD

- **File Systems Fated for Senescence? Nonsense, Says Science!** *Alex Conway, et al. FAST'17*

- SSD have 2 to 5 times slower read performance when accessing fragmented files

Fragmentation in SSD

- **SSD**

- **File Systems Fated for Senescence? Nonsense, Says Science!** *Alex Conway, et al. FAST'17*

- SSD have 2 to 5 times slower read performance when accessing fragmented files

- **FragPicker: A New Defragmentation Tool for Modern Storage Devices**

- Park, Jonggyu, and Young Ik Eom. ACM SIGOPS'21*

- Claims that SSD's performance degradation is mainly from request splitting

Fragmentation in SSD

■ SSD

- **File Systems Fated for Senescence? Nonsense, Says Science!** *Alex Conway, et al. FAST'17*

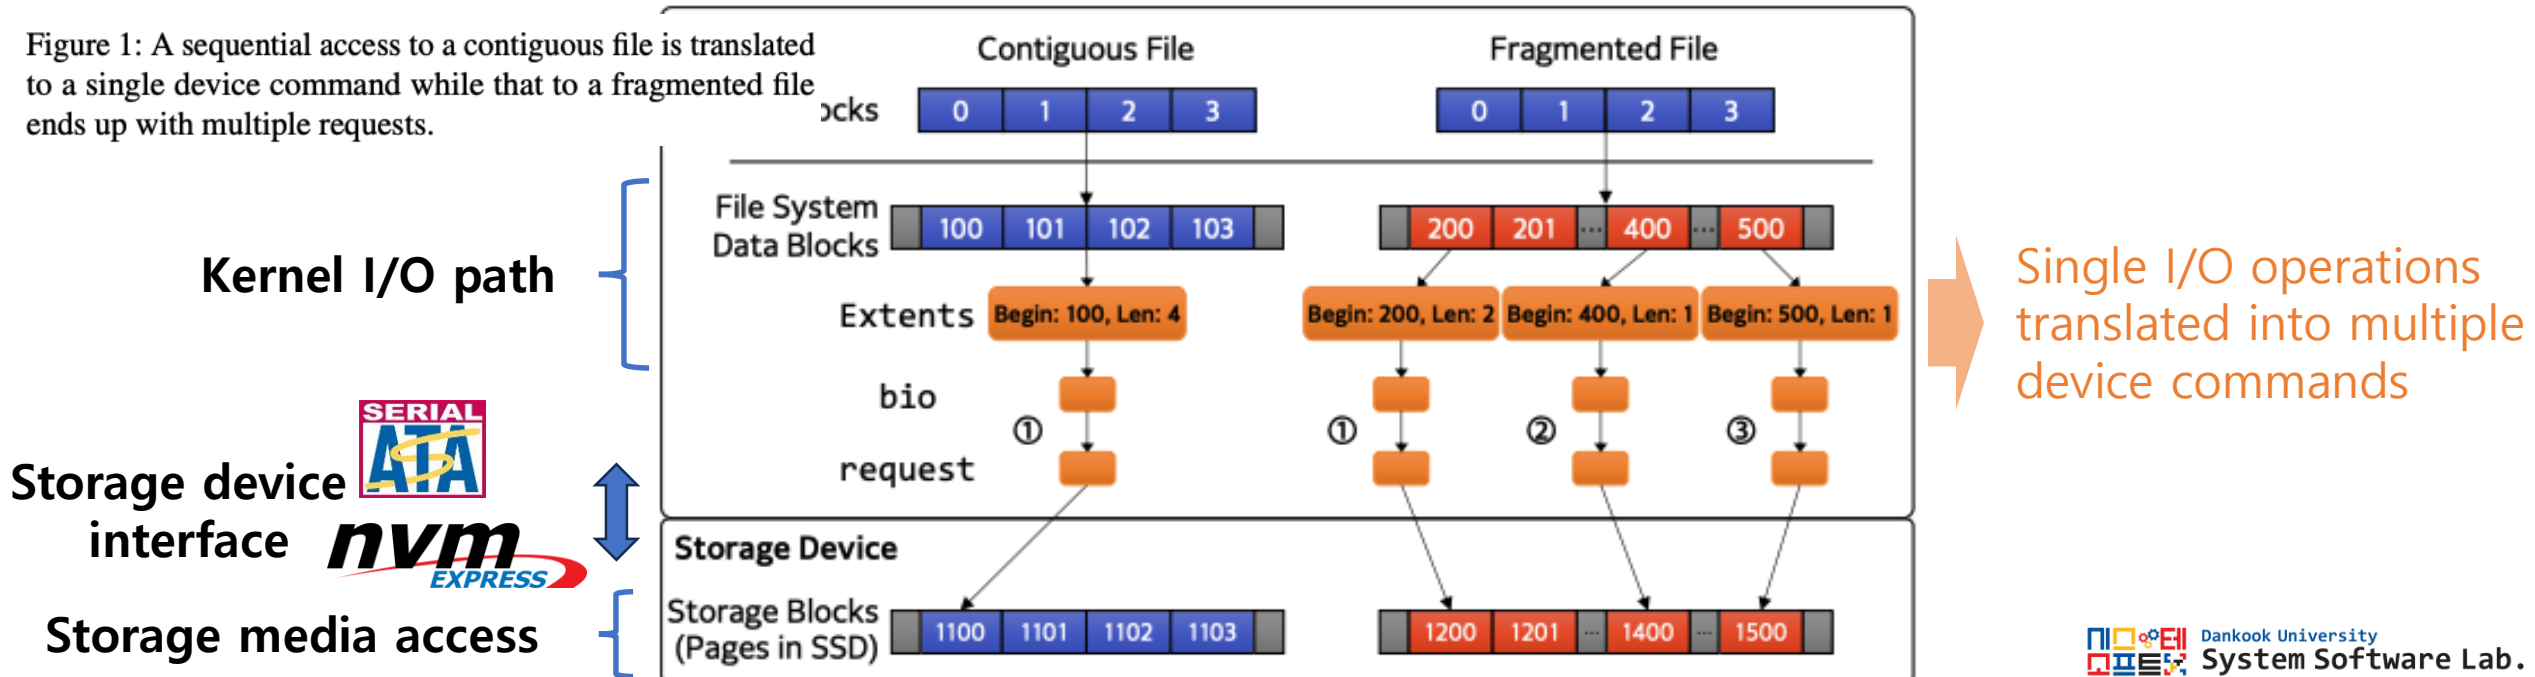
→ SSD have 2 to 5 times slower read performance when accessing fragmented files

- **FragPicker: A New Defragmentation Tool for Modern Storage Devices**

Park, Jonggyu, and Young Ik Eom. ACM SIGOPS'21


→ Claims that SSD's performance degradation is mainly from request splitting

Figure 1: A sequential access to a contiguous file is translated to a single device command while that to a fragmented file ends up with multiple requests.

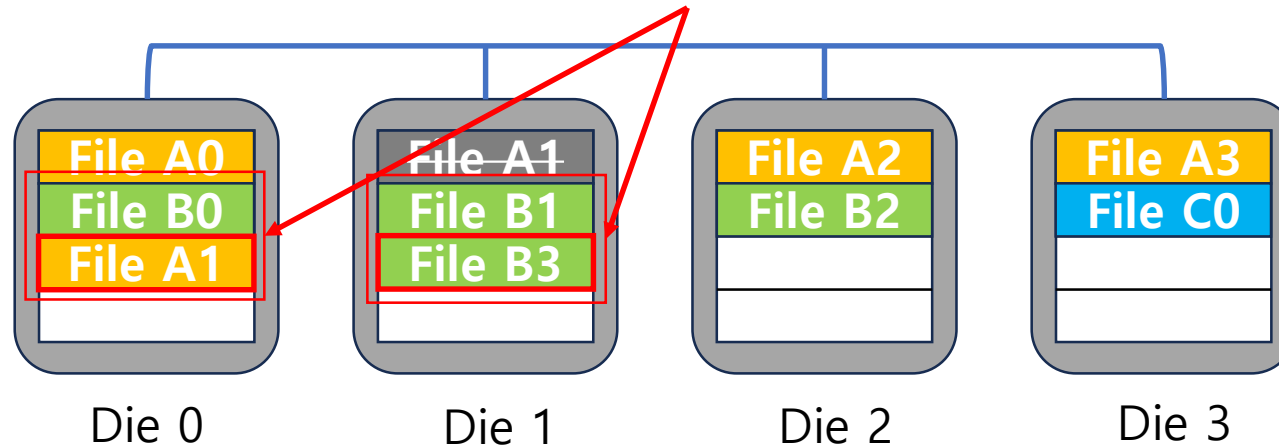


Fragmentation in SSD

■ Issue



Die-level collision



- File A Overwrite <A1>
- File B Append <B3>

Assigns in **round-robin manner**

Fragmentation in SSD

- Reason why?

→ Die Level Collisions!


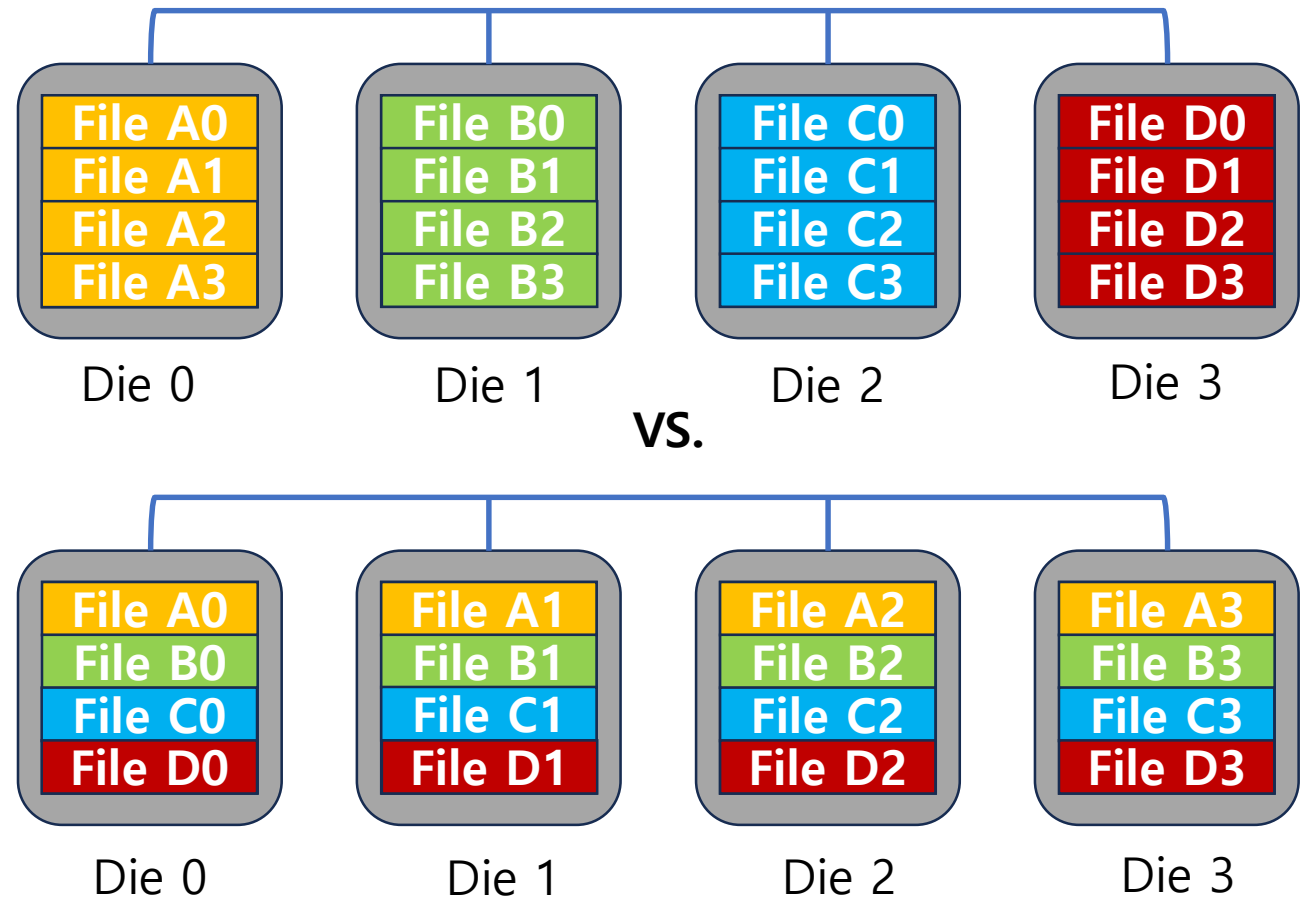


Figure 2: Data placement of three files in a flash SSD where one is contiguous and the other two are fragmented.

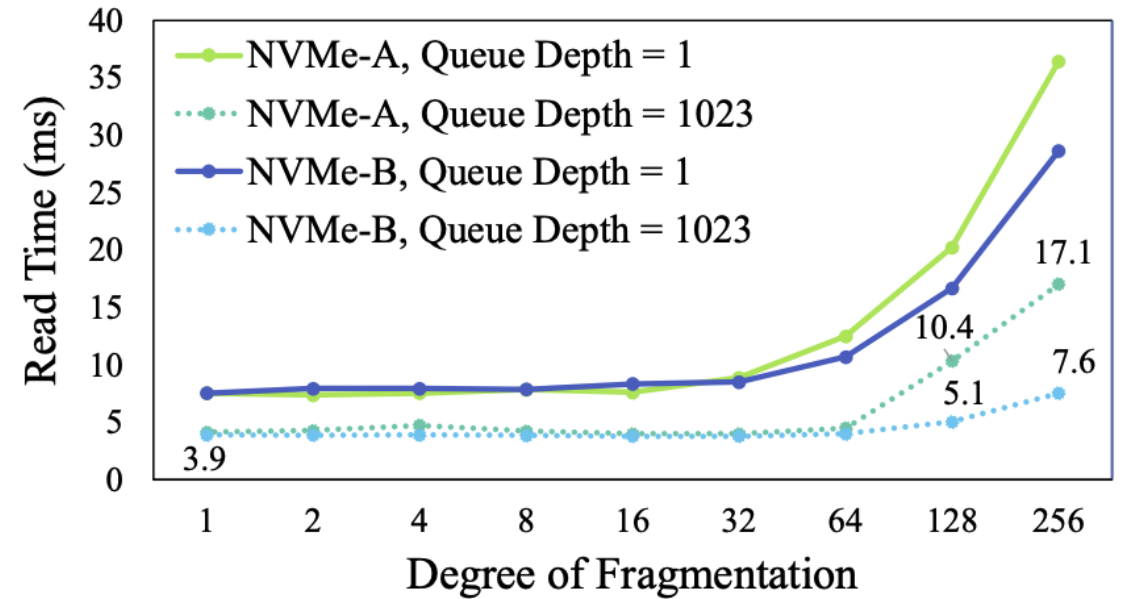
→ Die can only process one request at a time



Analysis of File Fragmentation

Table 1: System configurations for experiments.

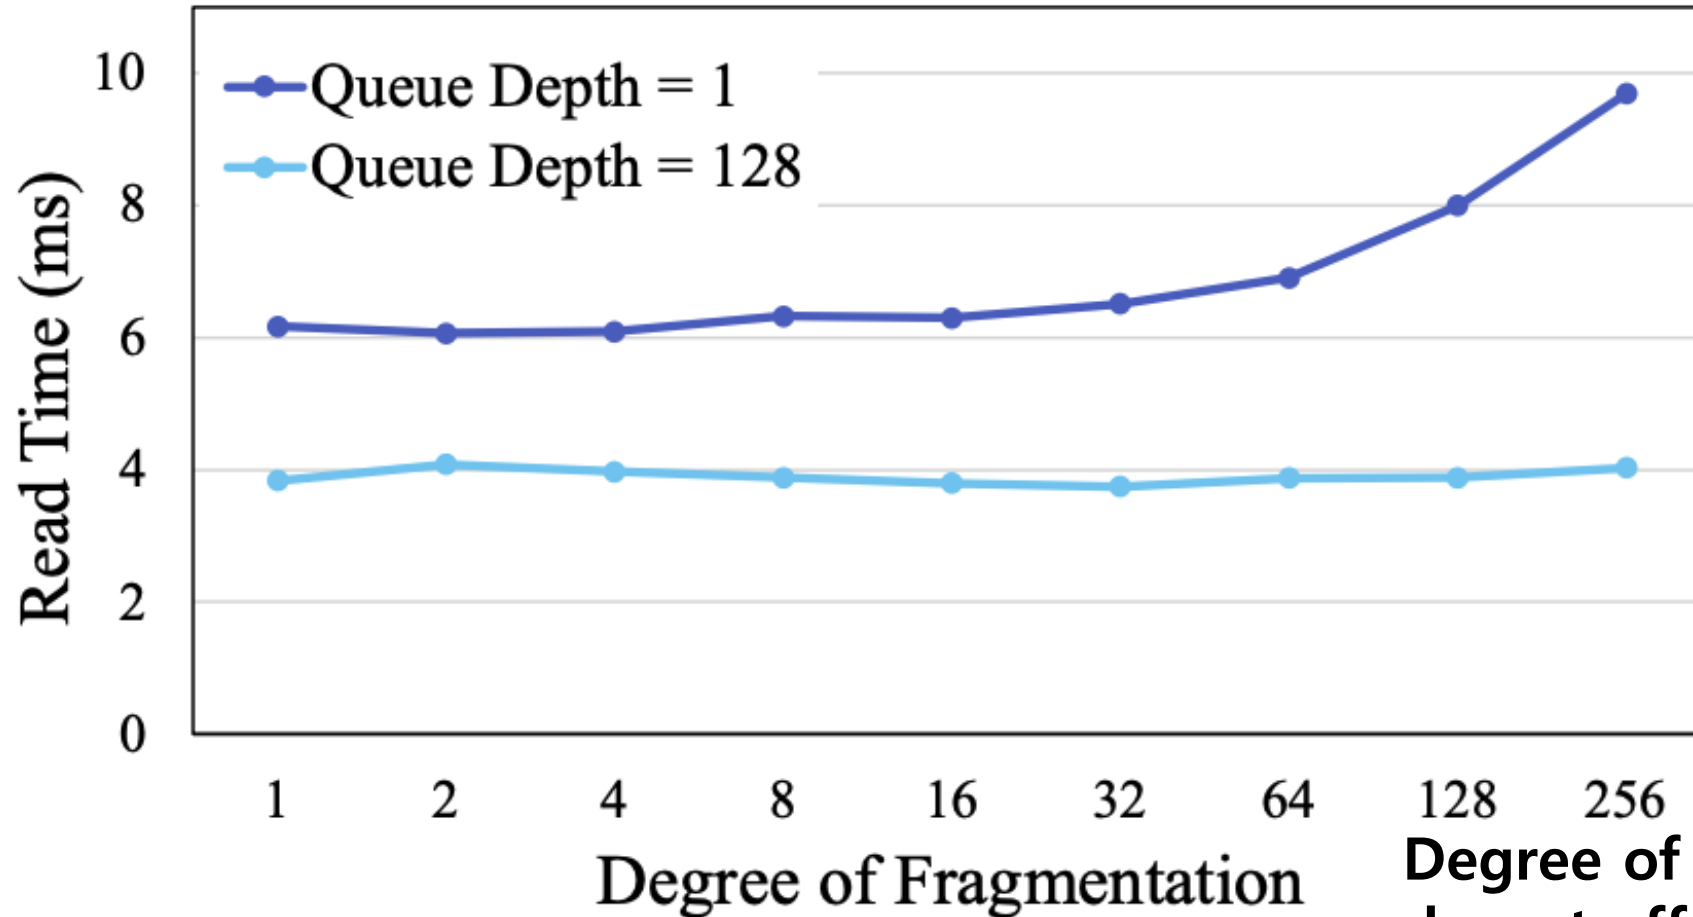
Processor	Intel Xeon Gold 6138 2.0 GHz, 160-Core
Chipset	Intel C621
Memory	DDR4 2666 MHz, 32 GB x16
OS	Ubuntu 20.04 Server (kernel v5.15.0)
Interface	PCIe Gen 3 x4 and SATA 3.0
Storage	NVMe-A: Samsung 980 PRO 1 TB
	NVMe-B: WD Black SN850 1 TB
	NVMe-C: SK Hynix Platinum P41 1 TB
	NVMe-D: Crucial P5 Plus 1 TB
	SATA-A: Samsung 870 EVO 500 GB
	SATA-B: WD Blue SA510 500 GB



Degree of Fragmentation causes performance degradation.

Analysis of File Fragmentation

■ RAMDisk



Degree of Fragmentation
do not affect to read time in ramdisk
→ No impact from request splitting

Analysis of File Fragmentation

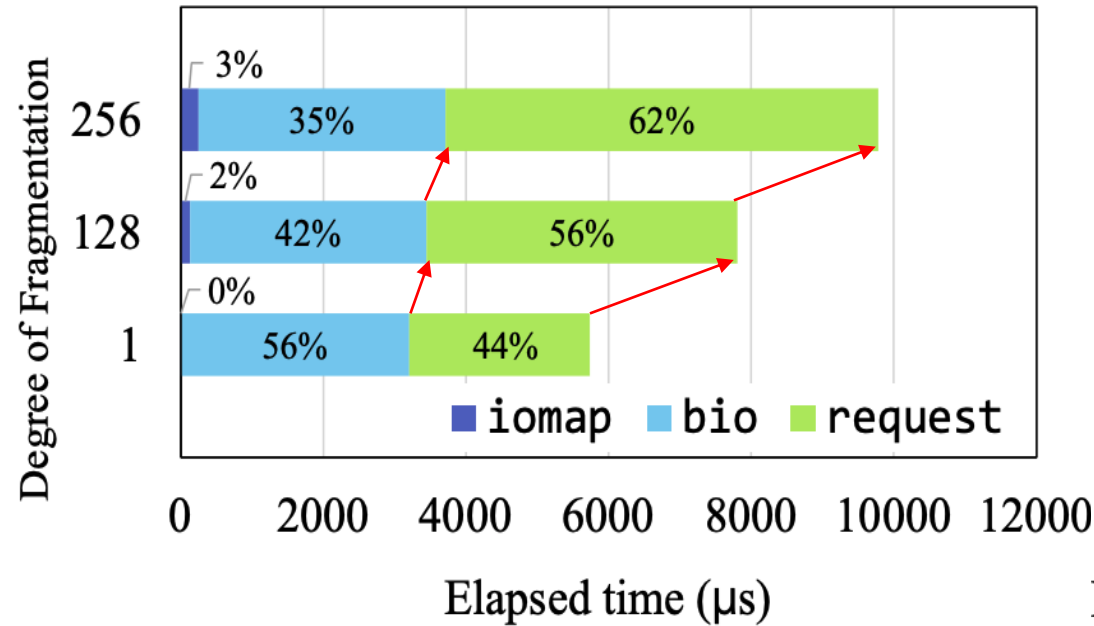


Figure 5: Time composition for creating request data structures in the kernel I/O path depending on File's DoF.

**Request time increased proportionally
With the increase in the
Degree of Fragmentation(DoF)**





Figure 6: Reduction of read time due to the overlap of storage operations and request creation when File's DoF is 128.

**“Kernel I/O path can be masked
by I/O queueing”**

Analysis of File Fragmentation




Result

→ Request splitting overhead in the kernel I/O path is negligible


→ Request splitting overhead is mitigated when issuing I/O operations asynchronously through command queueing

Analysis of File Fragmentation

- Page misalignment



Write in round-robin manner



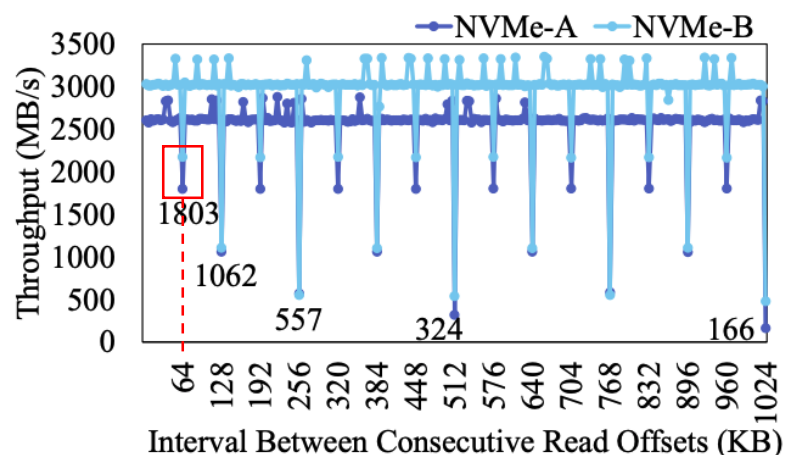
File A: up to 4x more bandwidth than single die

File B and C: 2x slower than File A

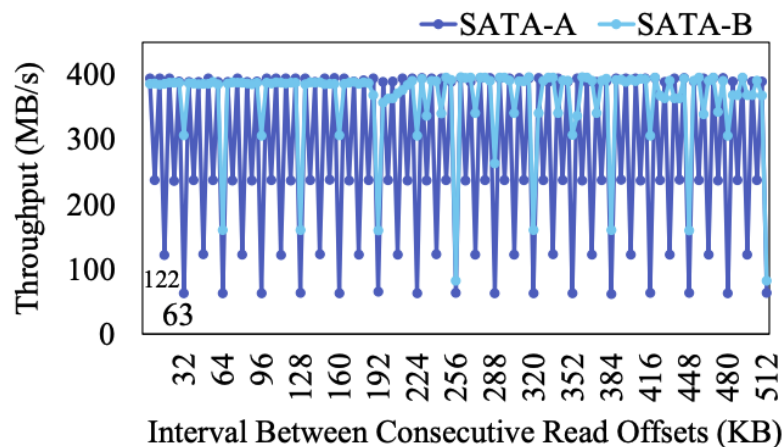
Die-level collision

Analysis of File Fragmentation

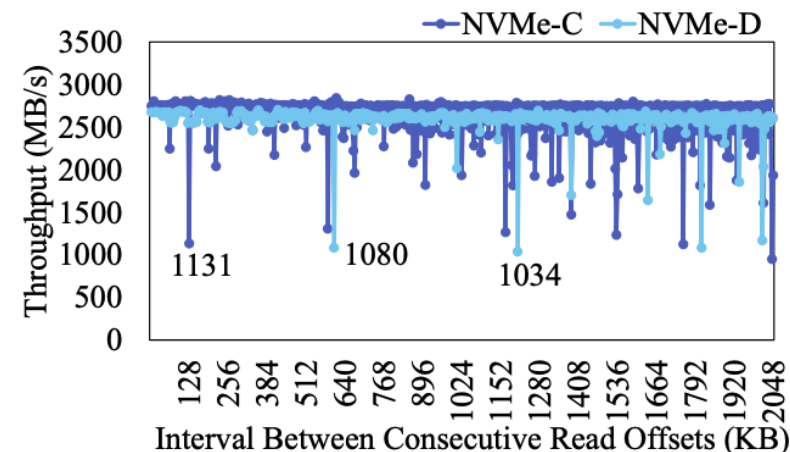
- Page misalignment



(a) NVMe-A/B



(b) SATA



(c) NVMe-C/D

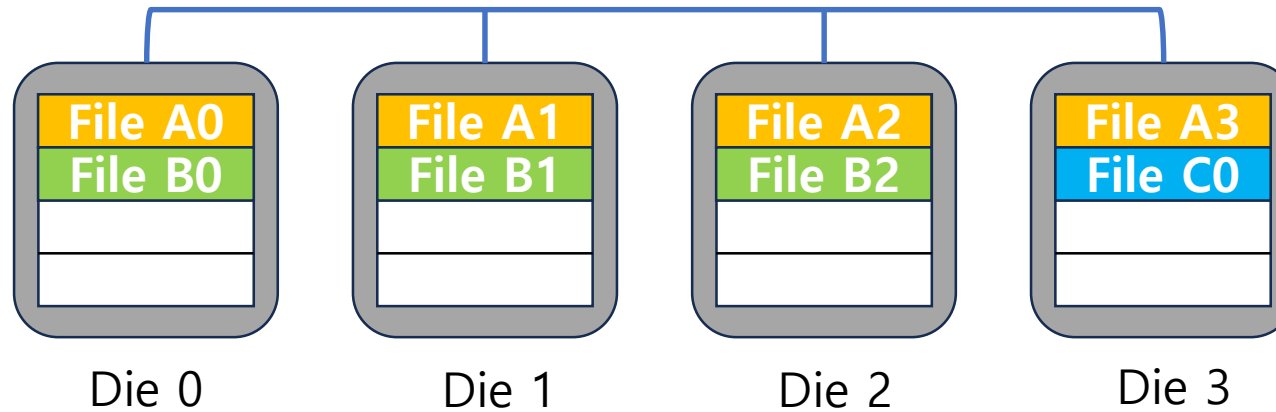
Figure 8: Throughput while varying the interval between starting points of consecutive read operations.

(a) Both NVMe SSD's page size is 16KB -> Allocates 2 pages per die

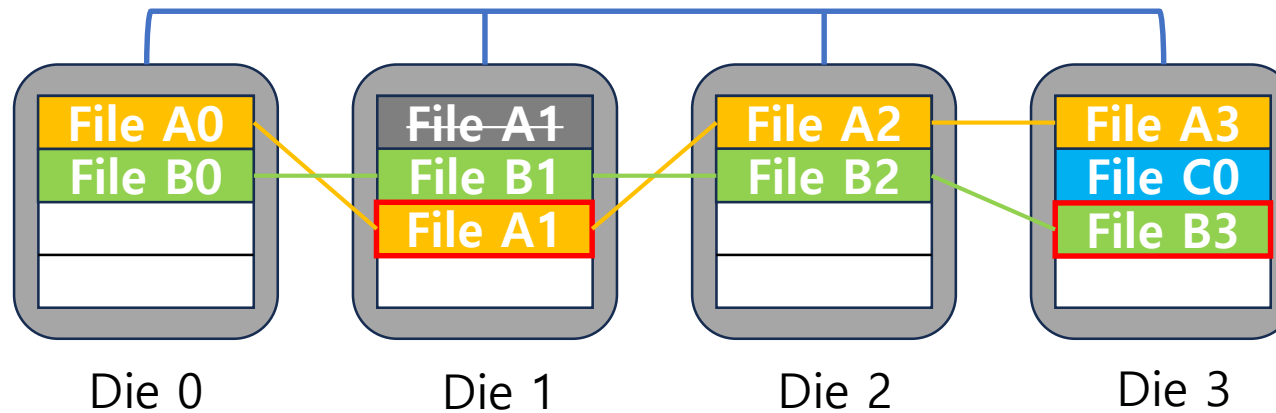
In SSD, file fragmentation leads to additional die-level collisions

Approach

- Using the given approach



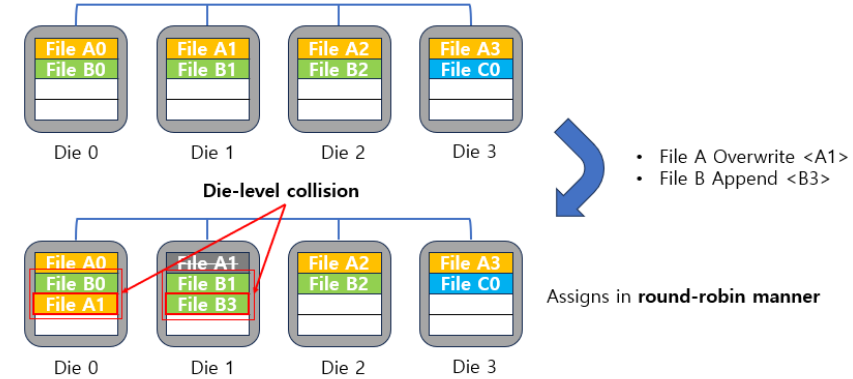
Locate in same die



Locate in subsequent die

Fragmentation in SSD

- Issue

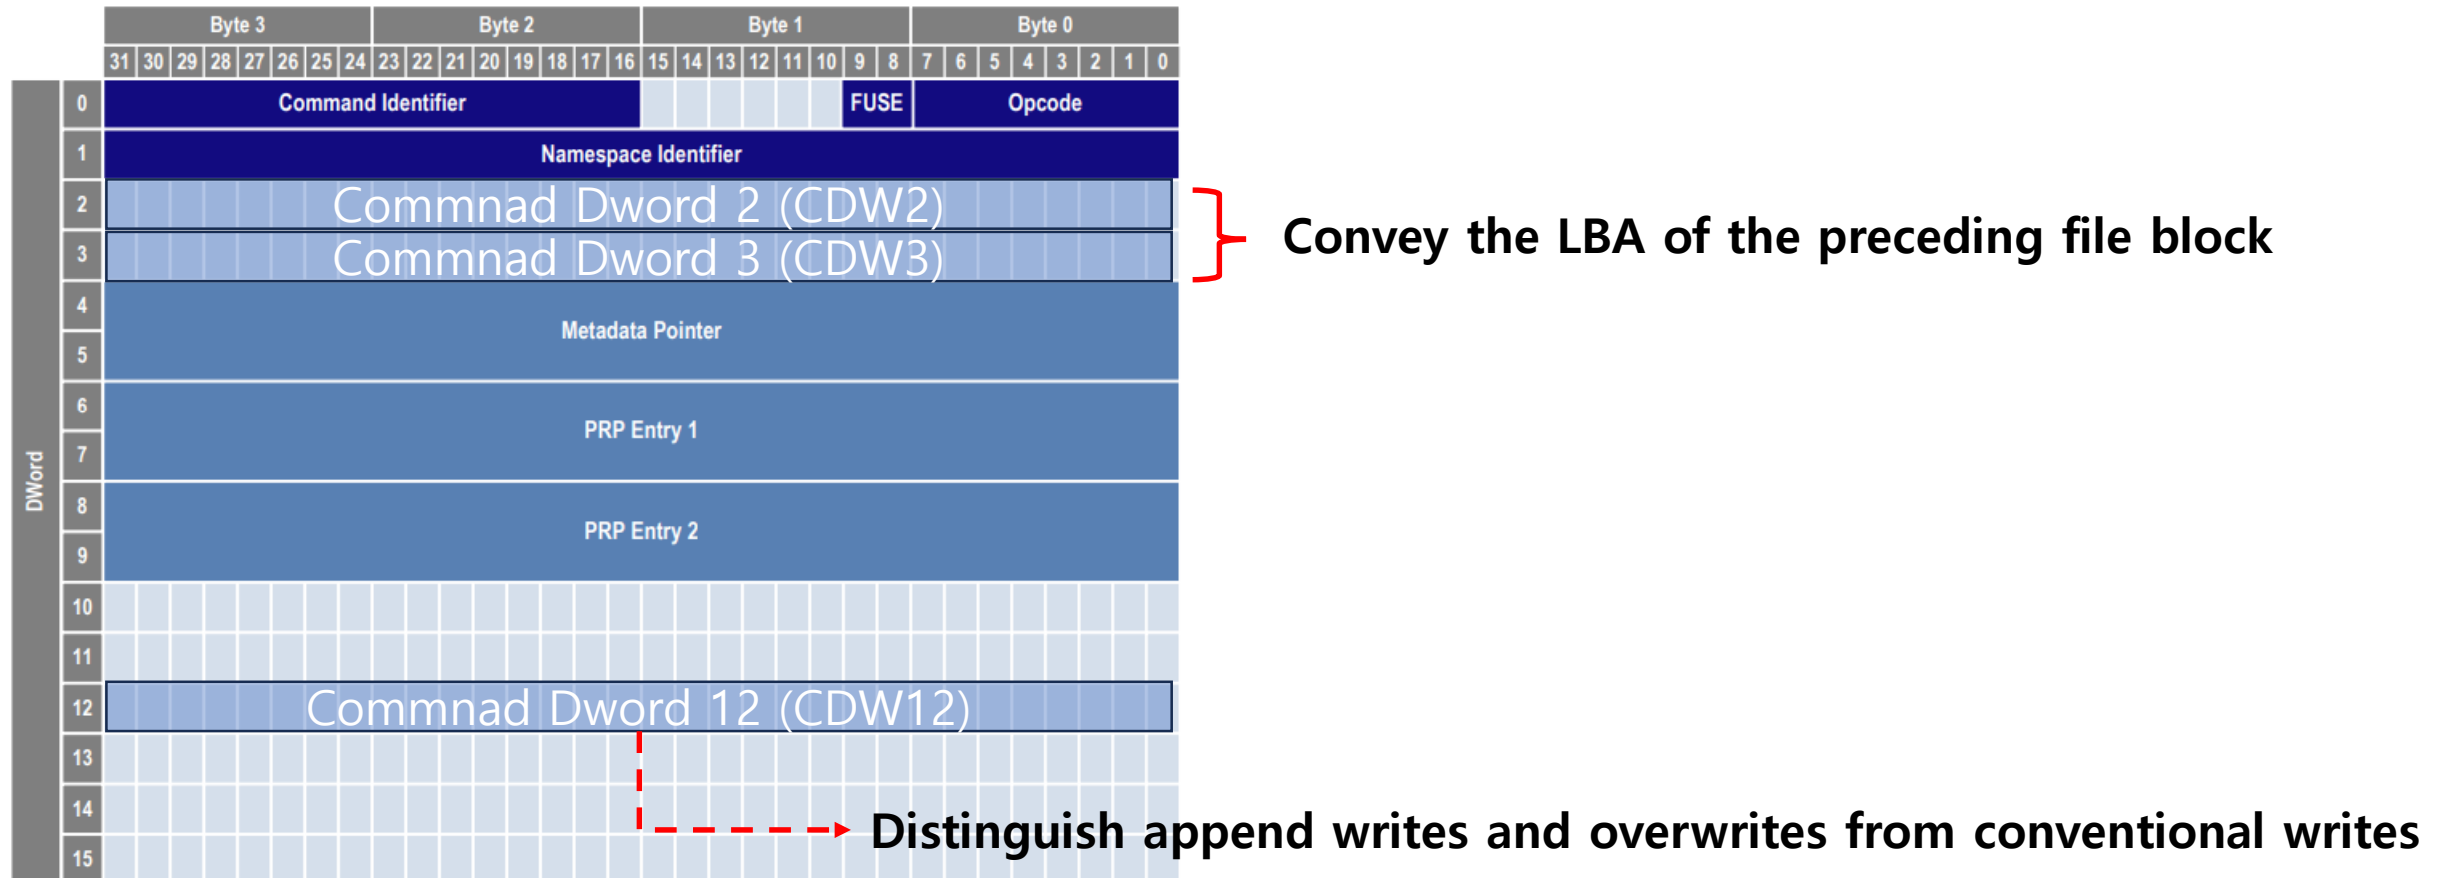


- File A Overwrite <A1>
- File B Append <B3>

Good parallelism!

Approach


- Uses NVMe protocol's write command
Submission Queue Entry (64B)




Evaluation

- Modified write patterns
- Showing read throughput

- Form a file by append 256 segments
- Each segment size
→ SSD's die allocation granularity
- Total file size = 8MB



(a) Append Write




(b) Overwrite

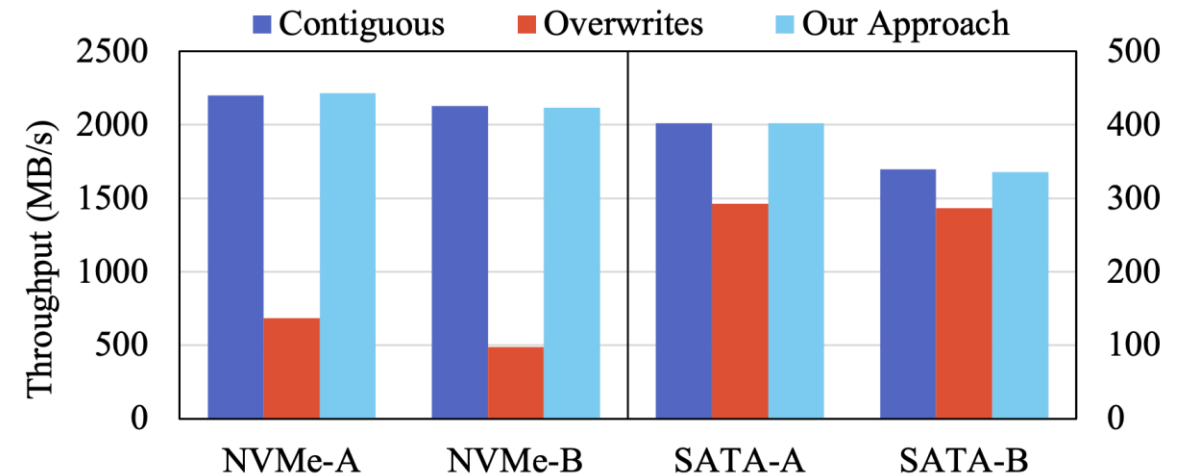
Evaluation

▪ Why does SATA SSDs performance degradation is less severe than NVMe?

- SATA3 Maximum throughput = **600MB/s**
- Smaller die allocation granularities in SATA SSD
- Adjusted final append's size to fit 8MB
- So only the initial segment of the file became fragmenated in SATA SSD



(a) Append Write



(b) Overwrite

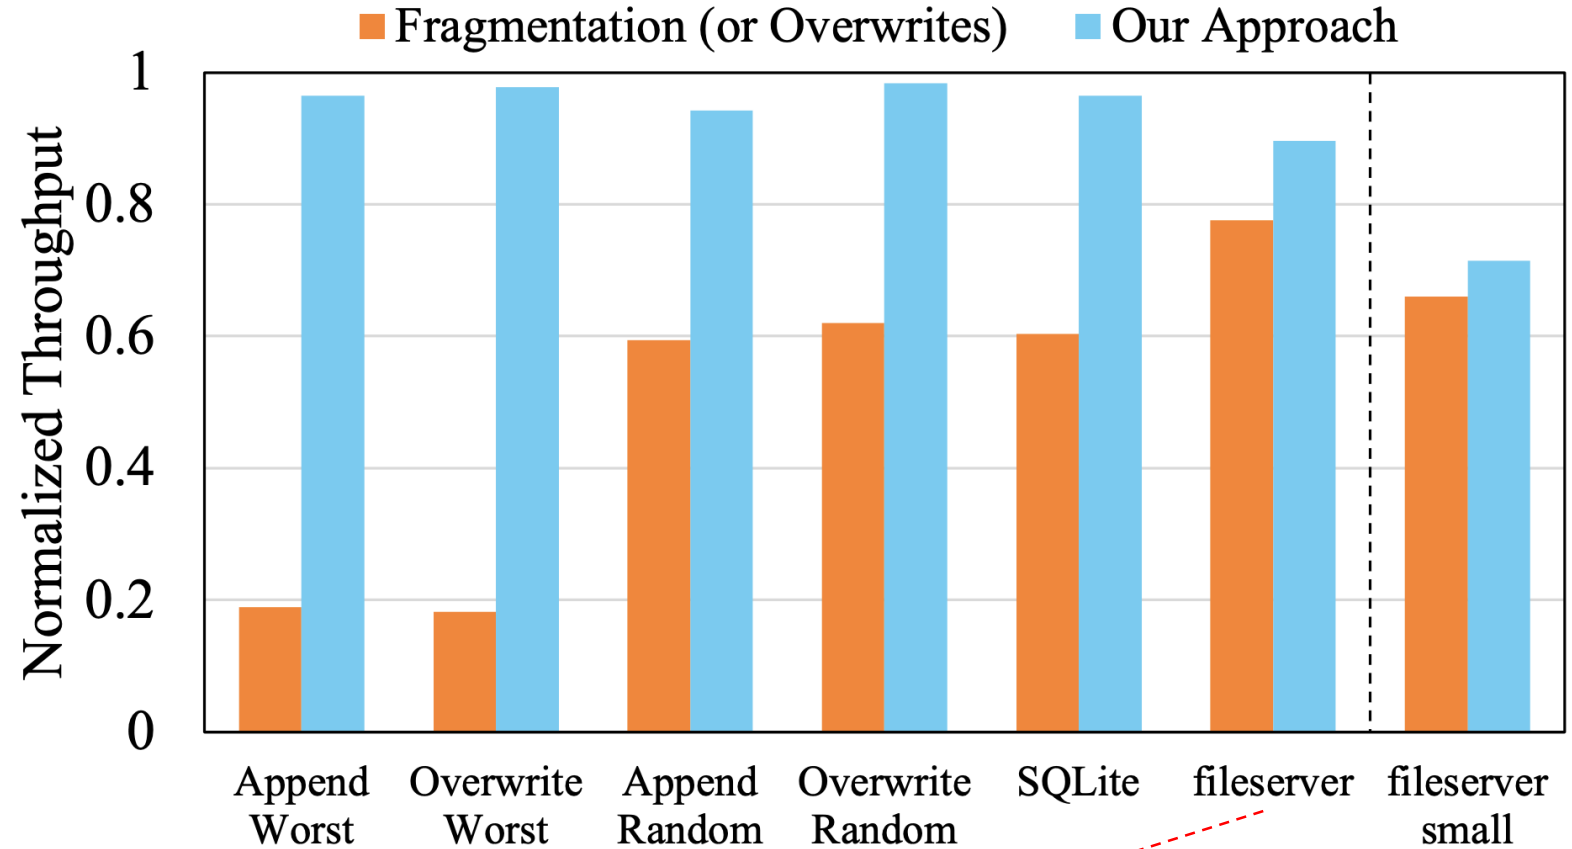
Evaluation

■ NVMeVirt

Table 2: Parameters used for NVMe emulation.

SSD	Capacity	60 GB
	Host Interface	PCIe Gen3 ×4
	FTL L2P Mapping	Page Mapping [1, 6]
	Channel Count	4
	Dies per Channel	2
Flash Memory [22]	Read/Write Unit Size	32 KB
	Read Time	36 μs
	Write Time	185 μs
	Channel Speed	800 Mbps

Mirrors the settings of NVMe-B



10 threads, 32KB size append writes

Worst case: located in single die

Reduced to 16KB

Conclusion

- **File fragmentation can indeed declines in read performance in SSD**
 - Because of die-level collisions rather than request splitting
 - Misalignments also happens when files are overwritten
- **Proposed NVMe command extension for better die-level parallelism**
 - Provide hints to SSD -> prevent additional die-level collisions caused by both file fragmentation and overwrites
 - Effectively suppresses the read performance degradation

Thank You !

2024. 08. 14

Presented by Juhyun Kim & Yongmin Lee

(email) , nascarf16@dankook.ac.kr