

Baleen: ML Admission & Prefetching for Flash Caches

Hao Wu[†], Carson Molder[§], Sathya Gunasekar[†], Jimmy Lu[†], Snehal Khandkar[†]Abhinav Sharma[†], Daniel S. Berger[‡], Nathan Beckmann, Gregory R. Ganger[†] Meta, [‡] Microsoft/UW, [§] UT Austin

FAST' 24

Source: American Museum of Natural History


2024.08.28

Presentation by Boseung Kim, Yejin Oh

bskim1102@dankook.ac.kr, yejinoh@dankook.ac.kr


Large-scale storage

- Large-scale storage uses HDD because of its cost-efficiency



Large-scale storage

- Large-scale storage uses HDD because of its cost-efficiency
- However, HDDs have **low bandwidth** and IOPS



Large-scale storage

- Large-scale storage uses HDD because of its cost-efficiency
- However, HDDs have **low bandwidth** and IOPS





Figure 2: Disk-head Time (DT) for one IO. When a HDD performs an IO, the disk head **seeks** before it **reads** data. For tiny IOs, throughput is limited by *IOPS*; for large IOs, by *bandwidth*. DT encompasses both metrics and generalizes to variable-size IOs.


Flash cache

- So, **flash caches** are utilized to absorb a fraction of requests




Flash cache

- So, flash caches are utilized to absorb a fraction of requests
- Flash provides higher IOPS, but it **wears out** as it is written
 - Specially in caching workload



Flash cache

- These SSD disadvantages require a proper **admission policy**
→ ML policies for flash cache is proposed



Challenges in the flash caching

- Flash admission
 - **Write cost** is determined at caching time
- ML admission
 - **Correct optimization metric** not obvious → Disk-head Time (DT)
 - train on all accesses in a trace, but focus on those that **result in misses** → Episode
- Prefetching policy
 - **Misprefetching** is costly as it consumes writes and extra DT

Challenges in the flash caching

- Flash admission
 - Write cost is determined at caching time
- ML admission
 - Correct optimization metric not obvious → **Disk-head Time (DT)**
 - Train on all accesses in a trace, but focus on those that result in misses → Episode
- Prefetching policy
 - Misprefetching is costly as it consumes writes and extra DT

Disk-head time(DT)

- End-to-end throughput metric to evaluate admission
- Cost of serving requests to the backend
 - Seek time + Read time, $DT_i = t_{seek} + n \cdot t_{read}$
 - Backend load: total DT needed to serve misses

$$Util_{DT} = \frac{\sum_i DT_i}{DT_{Provisioned}}, \quad \sum_i DT_i = Fetches_{IOs} \cdot t_{seek} + Fetches_{Bytes} \cdot t_{read} \quad (1)$$

Disk-head time(DT)

- End-to-end throughput metric to evaluate admission
- Cost of serving requests to the backend
 - Seek time + Read time, $DT_i = t_{seek} + n \cdot t_{read}$
 - Backend load: total DT needed to serve misses

$$Util_{DT} = \frac{\sum_i DT_i}{DT_{Provisioned}}, \quad \sum_i DT_i = Fetches_{IOs} \cdot t_{seek} + Fetches_{Bytes} \cdot t_{read} \quad (1)$$




Figure 3: DT validated in production. Our DT formula (plugging counters into Eq 1) matches measured disk utilization (blue) closely. The peak of 58% occurs on Day 0.

Use DT to address sub-problems

- Admission
 - Compare saved disk-head time and write cost when caching items
- Prefetching
 - Trade off DT saved from hits and wasted from incorrect prefetches
- Eviction
 - Using simple LRU
 - Leaving ML-based eviction policies for future work...

Challenges in the flash caching


- Flash admission
 - Write cost is determined at caching time
- ML admission
 - Correct optimization metric not obvious → Disk-head Time (DT)
 - Train on all accesses in a trace, but focus on those that result in misses → **Episode**
- Prefetching policy
 - Misprefetching is costly as it consumes writes and extra DT

Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall


Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall




Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall




Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall



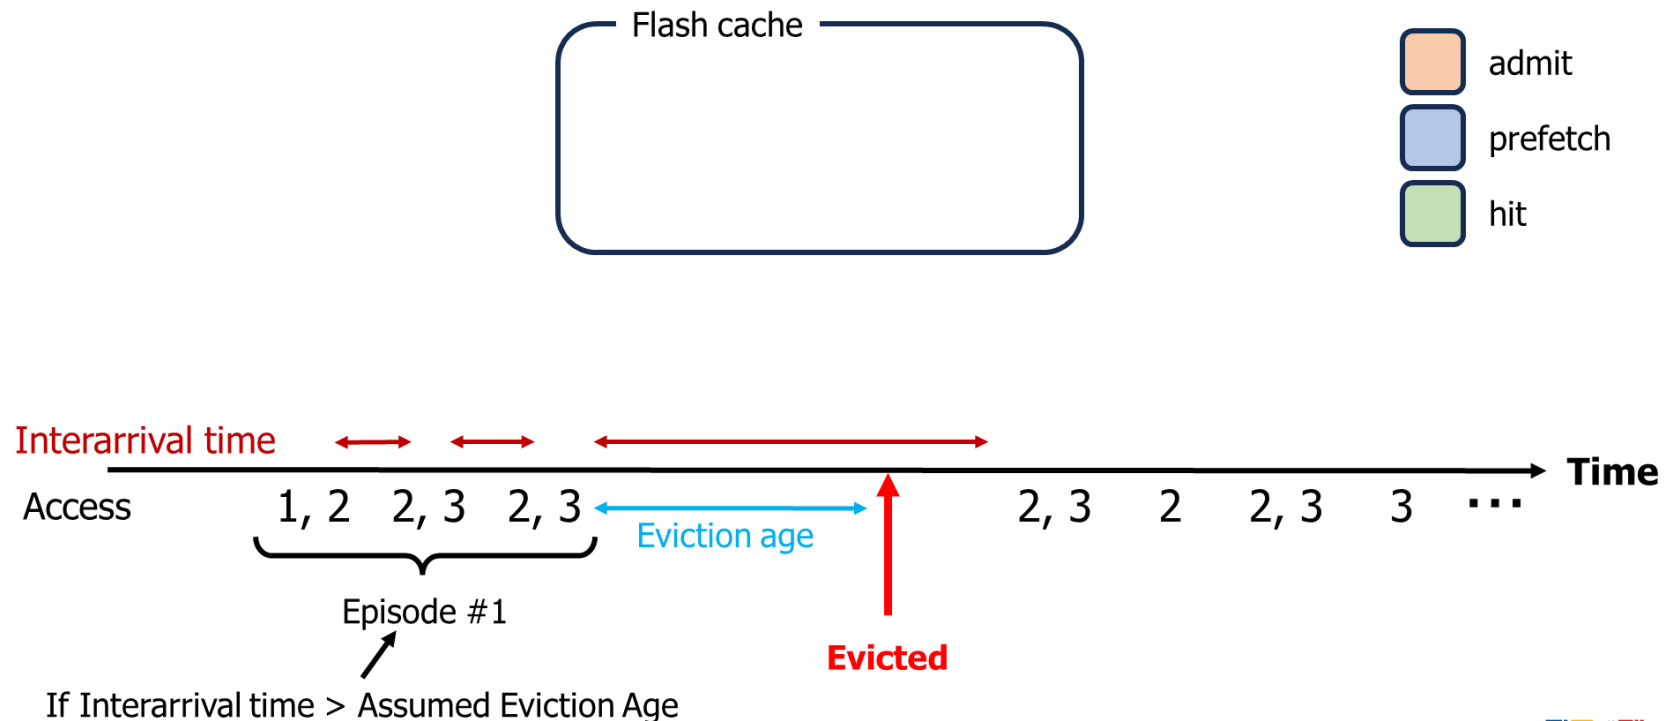
Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall



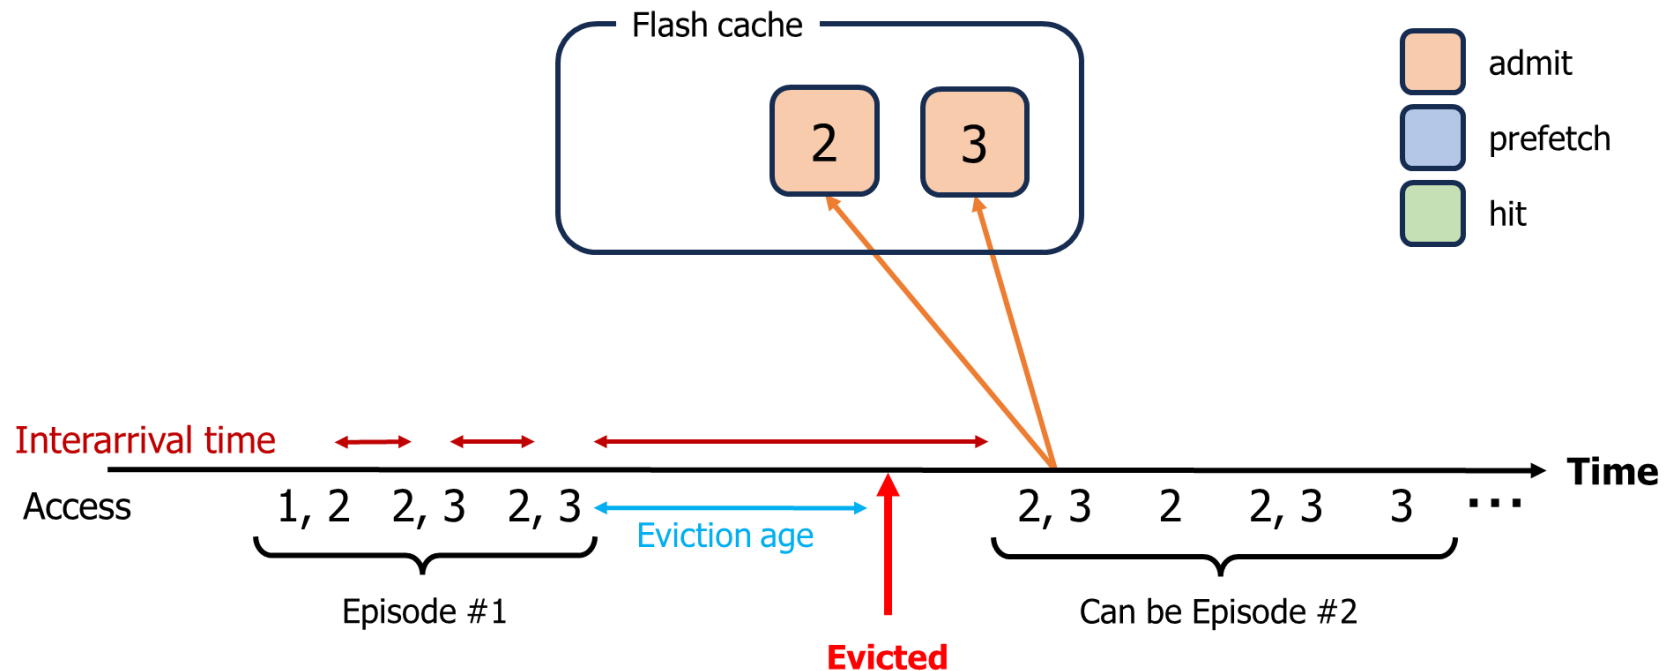
Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall



Episode

- Definition: a sequence of accesses that would be hits if the corresponding item was admitted
 - The entire process of an item entering and leaving the cache (=life span)
 - It allows us to consider how often and when an item is used in the cache overall




Episode's cost-benefit

- Episode's **starting point**
 - The earlier it starts, the more the write cost can be amortized
- Episode's **length**
 - The longer the episode, the higher the probability of a cache hit each time the item is used
- Episode's **access pattern**
 - The more accesses there are, the greater the benefit


Optimal AP using episode

- DT saved: The amount of data transfer saved by caching the episode
- Size: The amount of flash writes required to cache the episode
- $Score(Episode) = \frac{DT_{Saved}(Episode)}{Size(Episode)}$




Optimal AP using episode

- DT saved: The amount of data transfer saved by caching the episode
- Size: The amount of flash writes required to cache the episode
- $Score(Episode) = \frac{DTSaved(Episode)}{Size(Episode)}$



Optimal AP using episode

- $Score(EP) = \frac{DT\ Saved(Ep)}{FlashWrites(Ep)} = \frac{27ms}{3\ flash\ writes} = 9$
- $Score(Ep) > Cutoff_{Target\ Flash\ Write\ Rate}$



Challenges in the flash caching

- Flash admission
 - Write cost is determined at caching time
- ML admission
 - Correct optimization metric not obvious → Disk-head Time (DT)
 - Train on all accesses in a trace, but focus on those that result in misses → Episode
- Prefetching policy
 - Misprefetching is costly as it consumes writes and extra DT

Prefetching Policy

- Benefit: decreased DT
- Cost:
 - Overfetch: Minimal DT reduction effect
 - Underfetch: Excessive writes leading to performance degradation
- When?
 - OPT-Ep-Start: Fetch all necessary segments from the start of the episode to maximize DT reduction
- What?
 - Whole-Block: Prefetch the entire 8MB block
 - OPT-Range: Prefetch only the optimal segment range determined by the episode model

Baleen

- Provide episode-based solutions
 - How to train **ML-based admission policy**
 - Using **prefetch** to improve beyond admission policies

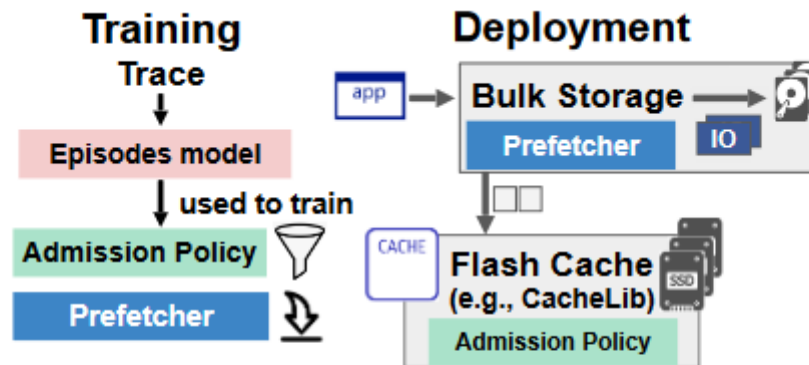


Figure 6: Architecture. An admission policy in CacheLib decides whether to admit items into flash. Prefetching (preloading of data beyond current request) takes place in Tectonic.

Baleen – ML admission policy

- Training Baleen's ML admission policy
 - 1) How to generate training data and labels from episodes
 - 2) What features and architecture we use for the ML admission model
 - 3) how we determine appropriate values for training parameters
 - 4) how we implement ML admission in CacheLib

Baleen – ML admission policy

- Training Baleen's ML admission policy
 - 1) how we generate training data and labels from episodes
 - Only the first 6 accesses from each episode are incorporated into training data
 - Baleen learns to imitate OPT

Baleen – ML admission policy

- Training Baleen's ML admission policy

2) what **features** and architecture we use for the ML admission model

- Metadata feature
 - **Request Source Identification:** identify the provenance of request (namespace, user)
- Online dynamic feature
 - **Access Count Tracking:** Records the number of times each block is accessed
 - **IO Statistics:** Tracks the number of input/output operations for each block and the cumulative segment IO operations.

Baleen – ML admission policy

- Training Baleen's ML admission policy
 - 2) what features and **architecture** we use for the ML admission model
 - Model: **binary classification**
 - the model outputs a probability for admitting misses
 - We admit misses if this probability exceeds the policy threshold

Baleen – ML admission policy

- Training Baleen's ML admission policy

3) how we determine appropriate values for training parameters

- Repeated Calculations to Determine Eviction Age and Policy Threshold Parameters
- Online simulation
 - Iterate until the measured average eviction age (EA) converges with the assumed EA
 - Calculate the policy threshold needed to achieve the desired flash write rate

4) how we implement ML admission in CacheLib

Baleen – ML Prefetcher

- ML-Range
 - Use of Two Regression Models
 1. Predict the **starting point** of an episode
 2. Predict the **ending point** of an episode
 - Train with size-related features
 - Access start index, access end index, access size

Baleen – ML Prefetcher

- ML-When
 - To reduce the effect of prefetching on average episode eviction age
 - Model: binary classification
 - Selected by ϵ

$$\text{ML-When}(eps) = PFBenefit_{eps}^{ML-Range} > \epsilon$$

Evaluation

- Comparative group
 - CoinFlip
 - On a miss, segments for an access are either all admitted, or not at all, with probability p
 - RejectX
 - Rejects a segment the first X times it is seen, use $X=1$ here
 - Past accesses are tracked using probabilistic data structures like Bloom filters
 - CacheLib with ML policy
 - Improved limitations of Flashield and non-episode-related features

Evaluation

- How to balance #HDD against # SSDs
 - TCO: total cost of HDD reads and written flash byte

$$\text{TCO}_1 \propto \underbrace{\frac{\text{PeakDT}_1}{\text{PeakDT}_0} \cdot \#HDDs_0}_{\text{Cost of HDD}} + \underbrace{\frac{\text{Cost}_{SSD}}{\text{Cost}_{HDD}} \cdot \frac{\text{FlashWR}_1}{\text{FlashWR}_0} \cdot \#SSDs_0}_{\text{Cost of SSD}} \quad (2)$$

- Lower peak load → Fewer hard disks → Lower TCO

Evaluation

TCO: total cost of HDD reads and written flash byte

$$TCO \propto Cost_{HDD} \cdot \#HDDs + Cost_{SSD} \cdot \#SSDs$$

- Baleen reduce TCO and Peak DT




Figure 8: Baleen-TCO reduces TCO.




Figure 9: Baleen reduces Peak DT.

Training on episodes is effective for ML prefetching, and allowing Baleen to efficiently reduce TCO and peak DT.

Evaluation

- ML-guided prefetch




Figure 13: ML-Range saves Peak DT. ML-Range outperforms the baseline (whole block) and No Prefetching at the expense of Median DT.




Figure 14: Choose *when* to prefetch. Indiscriminate prefetching (on Every Miss) can hurt. Using ML-When or Partial Hit reduces Peak DT without compromising Median DT.

ML-Range efficiently reduce peak DT by limiting the number of segments prefetched

ML-When helps Baleen discriminate between beneficial and bad prefetching

Evaluation

- Explicitly optimizing Peak DT




Figure 15: Testbed backend load on Region1. Day 1 (shaded) is used as training data. Peak is on Day 5 and is lowest for Baleen.

Baleen can be aware of the current load level and able to adapt to optimize Peak DT

Conclusion

- Large-scale servers use Flash Cache to mitigate the slow bandwidth of HDDs
 - However, due to the limited lifespan of SSDs, proper admission policies are required
- Instead of using traditional performance evaluation metric, which are inadequate, Baleen used DT and TCO as performance measures
- Baleen used an episode model to improve ML learning efficiency.
 - ML-range: predict a range of segments for prefetching
 - ML-When: Avoid waste caused by misprefetching
- As a result, Baleen lowered DT by 12% and TCO by 17%, effectively reducing costs.

Thank you

Q & A

Optimal AP using episode

- 왜 optimal 한가?

- 전체 접근 이력을 사용하여 각 episode에 대한 완전한 정보를 가지고 있다
- 따라서 장기적인 이득이 극대화된다.
- Budget 고려함
- 이론적으로 최적 결과가 된다.

Workload

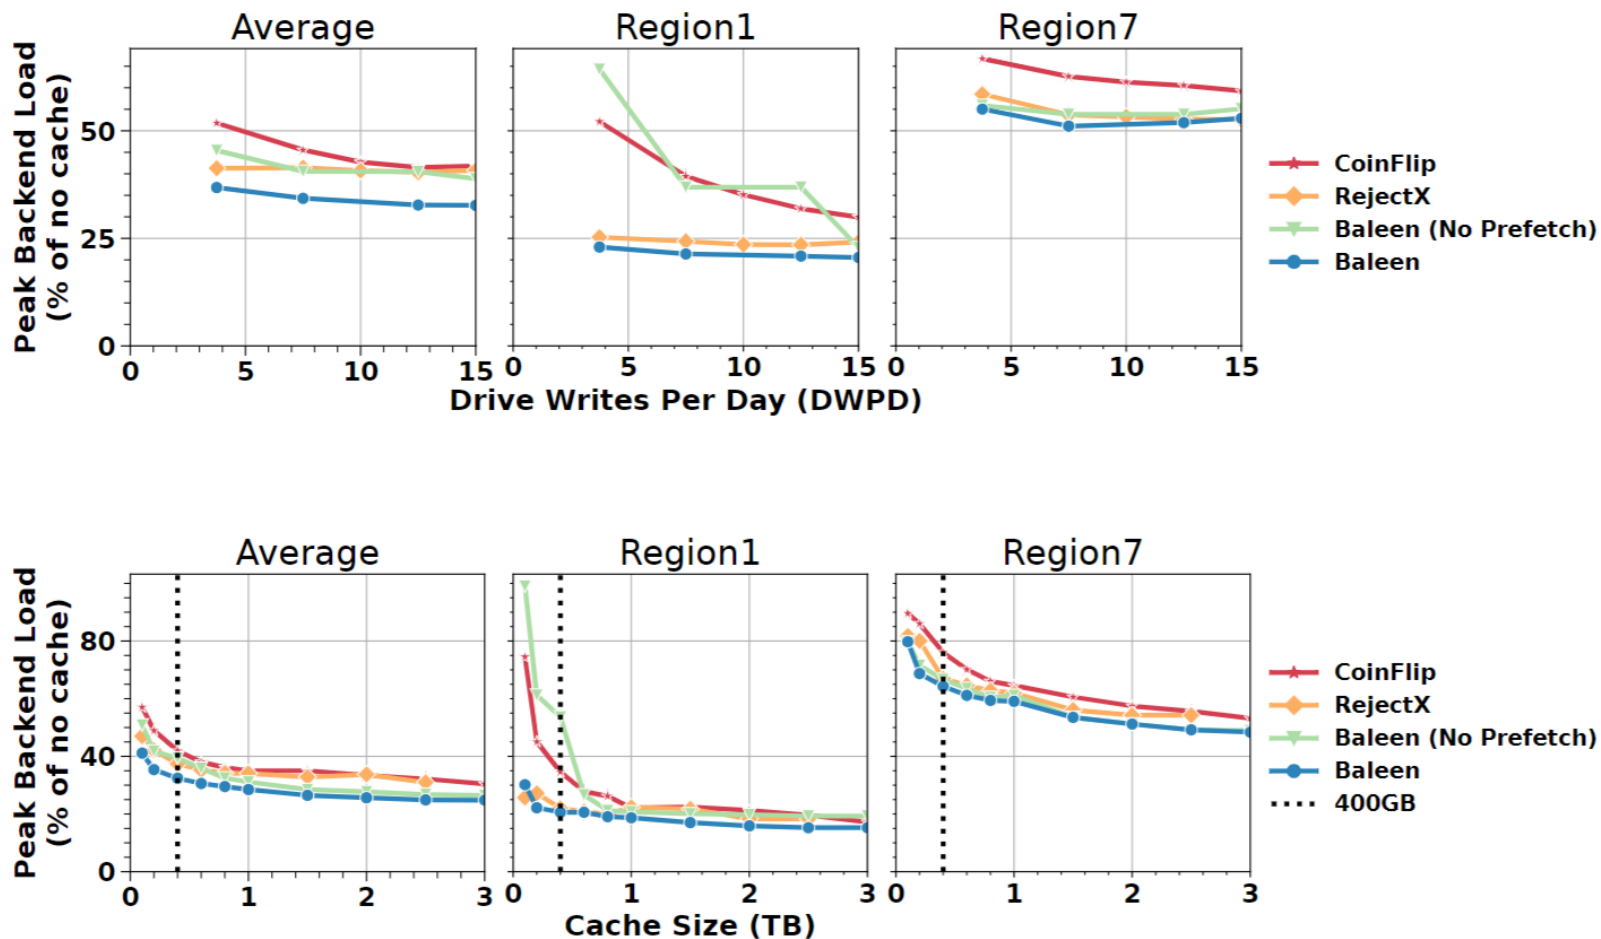


Figure 10: Benefits at higher write rates & cache sizes.